DQN-LSTM for Event Summarization

Chris Kedzie

kedzie@cs.columbia.edu

Francisco Javier Arceo

fja2114@columbia.edu

October 8, 2016

Abstract

I will insert something here[1]

1 Background

2 Extractive Streaming Summarization

In the extractive streaming summarization task, we are given as input a query, e.g. a short text description of a topic or an event, and document stream, a time ordered set of sentences relevant to the query. Starting with an initially empty summary, an extractive, streaming summarization algorithm is intended to examine each sentence in order and when new and important (relative to the query) information is identified, add that sentence to the summary.

Implicit to this problem is the notion of system time – the summarization algorithm can only examine sentences that occur in the stream before the current system time. Advancing the system time gives the algorithm access to more sentences, although in practice the stream is sufficiently large enough that choices have to be made about how much history can be kept in memory. For many domains, e.g. crisis informatics, it is preferable for a summarization algorithm to identify important information as early as possible, and so the objective function should penalize a large disparity between the time a piece of information is first available to the algorithm and the system time at which that information is actually added to the summary.

Previous work in this area has either incremented the system time in fixed increments (e.g. an hour) [5, 4] or operated in a fully online setting [2, 3]. In both

---

[1] https://github.com/franciscojavierarceo/DQN-Event-Summarization

1

cases explicitly modeling the current state of the summary, the stream, and their relationship with the query is quite complicated and exhibits non-linear dynamics that are difficult characterize in traditional feature based models.

Additionaly, the structured nature of the sentence selection task (sentence selection is highly dependent on the current summary state) suggests that imitation or reinforcement learning are necessary to obtain parity between training and testing feature distributions.

This leads us to explore deep Q networks (DQN) for two reasons. First, both the representation and mode of interaction between the stream, summary, and query can be learned. Second, the learned Q function (plus random noise) controls the state space that is explored, ensuring more consistancy between train and test distribtuions than for example naive imitation learning (possibly – I'm not totally happy with this sentence).

2.1 Algorithm

2.2 Problem Definition

DQN learns by using an a ϵ-greedy search policy to generate a sequence of state, action, next state, reward 4-tuples using the current learned Q-network to evaluate candidate actions. These tuples are sampled from and used to estimate the true Q function.

States In our setup a state $s(X_{\leq t}, \tilde{Y}_{\leq t}, q)$ is a function of the stream X observed up to the current system time t, the state of the current summary $\tilde{Y}$ at system time t, and the query q. For brevity we will use $s(t, q)$ where the dependence on $X_{\leq t}, \tilde{Y}_{\leq t}$ is assumed. $s(t, q)$ is itself three recurrent neural networks, one for encoding the summary, the stream, and the query respectively.

Actions The set of possible actions at each time step is $\mathcal{A} = \{select, skip\}$ where $select$ corresponds to adding the current sentence x_t to the summary and incrementing the current system time,

$$\tilde{Y}_{\leq t+1} = \tilde{Y}_{\leq t} \cup \{x_t\}$$
$$t = t + 1$$

or $skip$ where only t is incremented without changing the current summary

$$\tilde{Y}_{\leq t+1} = \tilde{Y}_{\leq t}$$
$$t = t + 1.$$

2

Algorithm 1 Streaming DQN-LSTM

Input: $\{x_q, \pi_q^*\}_{\forall q \in \mathcal{Q}}$, number of iterations N, mixture parameter ϵ
Output: $\hat{\pi}$

 1: **The notation is not very clear – WORK IN PROGRESS**
 2: Initialize experience memory Γ
 3: Initialize action-value function $\mathcal{Q}$ with random weights
 4: $i \leftarrow \emptyset$
 5: **for** $n \in \{0, 1, .., N-1\}$ **do**
 6:    **for** $q \in \mathcal{Q}$ **do**
 7:       **for** $t \in \{0, 1, ..., T-1\}$ **do**
 8:          Execute π_i t times and reach s_t
 9:          **for** $a \in \mathcal{A}(s_t)$ **do**
10:             With probability ϵ set $\pi^0 = \pi_q^*$, else π_i
11:             Compute $c_t(a)$ by executing π^0
12:             $\Gamma \leftarrow \Gamma \cup \{[\phi(s_t), a, c_t(a)]\}$
13:          **end for**
14:       **end for**
15:       Sample random minibatches of transitions γ_j from Γ
16:       Set $y_j = r_j$ for terminal ϕ_{j+1}
17:       Perform gradient step on $(y_j - Q(\phi_j, a_j; \theta))^2$
18:       $i \leftarrow i + 1$
19:    **end for**
20: **end for**

Algorithm 2 Streaming DQN-LSTM

Input: $\{x_q, \pi_q^*\}_{\forall q \in \mathcal{Q}}$, number of iterations N, mixture parameter ϵ
Output: $\hat{\pi}$

1: **The notation is not very clear – WORK IN PROGRESS**
2: Initialize π_0, Γ
3: Initialize action-value function $\mathcal{Q}$ with random weights
4: $i \leftarrow \emptyset$
5: **for** $n \in \{0, 1, .., N-1\}$ **do**
6:    **for** $q \in \mathcal{Q}$ **do**
7:      **for** $t \in \{0, 1, ..., T-1\}$ **do**
8:        Execute π_i t times and reach s_t
9:        **for** $a \in \mathcal{A}(s_t)$ **do**
10:          With probability ϵ set $\pi^0 = \pi_q^*$, else π_i
11:          Compute $c_t(a)$ by executing π^0
12:          $\Gamma \leftarrow \Gamma \cup \{[\phi(s_t), a, c_t(a)]\}$
13:        **end for**
14:      **end for**
15:      Sample random minibatches of transitions γ_j from Γ
16:      Set $y_j = r_j$ for terminal ϕ_{j+1}
17:      Perform gradient step on $(y_j - Q(\phi_j, a_j; \theta))^2$
18:      $i \leftarrow i + 1$
19:    **end for**
20: **end for**

Reward The reward for a given action will be measured by relative gain in ROUGE-2 F1 score of the predicted summary $\tilde{Y}_{\leq t}$ measured against a gold standard summary Y.

When only one gold summary reference is used, ROUGE-N Recall is calculated as

$$\text{ROUGE-NR}(\tilde{Y}, Y) = \frac{\sum_{g \in \text{ngrams}(Y,N)} \min\left(\text{count}(g, \tilde{Y}), \text{count}(g, Y)\right)}{\sum_{g \in \text{ngrams}(Y,N)} \text{count}(g, Y)}$$

where $\text{ngrams}(Y, N)$ returns the set of ngrams of order N in the summary Y and $\text{count}(g, Y)$ is the count of occurences of ngram g in Y.

Similarly, ROUGE-N Precision is calculated as

$$\text{ROUGE-NP}(\tilde{Y}, Y) = \frac{\sum_{g \in \text{ngrams}(Y,N)} \min\left(\text{count}(g, \tilde{Y}), \text{count}(g, Y)\right)}{\sum_{g \in \text{ngrams}(\tilde{Y},N)} \text{count}(g, \tilde{Y})}$$

and the F_1 is simply the harmonic mean of the two:

$$\text{ROUGE-NF1}(\tilde{Y}, Y) = \frac{2 \times \text{ROUGE-NP}(\tilde{Y}, Y) \times \text{ROUGE-NR}(\tilde{Y}, Y)}{\text{ROUGE-NP}(\tilde{Y}, Y) + \text{ROUGE-NR}(\tilde{Y}, Y)}$$

The reward r at time t is then:

$$r_t = \text{ROUGE-NF1}(\tilde{Y}_{\leq t+1}, Y) - \text{ROUGE-NF1}(\tilde{Y}_{\leq t}, Y)$$

TODO: Think about how to incorporate a time penalty into the reward.

Task	Status	Person
Read DQN Paper [7]	Complete	Francisco
Read Asynchronous RL Paper [6]	Briefed	Francisco
Read Regularization Paper [1]	Complete	Francisco
Read Disaster Summarization Paper [4]	Complete	Francisco
Read Sequential Decision Making Paper [3]	Complete	Francisco
Install Torch	Complete	Francisco
Small simulation of MLP in Torch	Complete	Francisco
Small simulation of DQN in Torch	Incomplete	Francisco
Download data	Incomplete	Francisco
Explore data	In-Progress	Francisco

DQN-LSTM for Event Summarization

References

[1] Minmin Chen, Zhixiang Xu, Kilian Weinberger, and Fei Sha. Marginalized denoising autoencoders for domain adaptation. *arXiv preprint arXiv:1206.4683*, 2012.

[2] Qi Guo, Fernando Diaz, and Elad Yom-Tov. Updating users about time critical events. In *European Conference on Information Retrieval*, pages 483–494. Springer, 2013.

[3] Chris Kedzie, Fernando Diaz, and Kathleen McKeown. Real-time web scale event summarization using sequential decision making. *arXiv preprint arXiv:1605.03664*, 2016.

[4] Chris Kedzie, Kathleen McKeown, and Fernando Diaz. Predicting salient updates for disaster summarization. In *Proceedings of the 53rd annual meeting of the ACL and the 7th international conference on natural language processing*, pages 1608–1617, 2015.

[5] Richard McCreadie, Craig Macdonald, and Iadh Ounis. Incremental update summarization: Adaptive sentence selection based on prevalence and novelty. In *Proceedings of the 23rd ACM International Conference on Conference on Information and Knowledge Management*, pages 301–310. ACM, 2014.

[6] Volodymyr Mnih, Adrià Puigdomènech Badia, Mehdi Mirza, Alex Graves, Timothy P. Lillicrap, Tim Harley, David Silver, and Koray Kavukcuoglu. Asynchronous methods for deep reinforcement learning. *CoRR*, abs/1602.01783, 2016.

[7] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Alex Graves, Ioannis Antonoglou, Daan Wierstra, and Martin A. Riedmiller. Playing atari with deep reinforcement learning. *CoRR*, abs/1312.5602, 2013.

[8] Karthik Narasimhan, Tejas D Kulkarni, and Regina Barzilay. Language understanding for textbased games using deep reinforcement learning. In *In Proceedings of the Conference on Empirical Methods in Natural Language Processing*. Citeseer, 2015.